

FreightIQ -- Agentic Document Verification

Part 2: applying Part 1's extraction + analytics to the SU → CG → Customer trade-document loop
GoComet Agentic AI PM Assignment · Part 2 (3-4 hour sprint) · Prepared by Kamal Gupta

Problem

SU submits shipment trade documents by email, and CG manually reads every field against customer-specific rules and types out amendment requests by hand, two to four cycles per shipment. Those rules live only in CG's heads, so a new hire is slow and error-prone for weeks, there is no audit trail if a dispute arises, and each amendment cycle adds four to twenty-four hours of delay with no visibility into how many documents are pending at any time.

Personas

Who	What they actually care about in this workflow
CG (Cargo / Control Group -- validator)	Catching every discrepancy before it reaches the customer -- a missed error becomes their problem (customs delay, contract penalty). Wants to stop reading every field by hand.
SU (Shipping Unit -- shipper's docs desk)	Getting the shipment cleared fast with minimal back-and-forth. A vague or slow amendment request costs them a full resubmission cycle and delays their own dispatch targets.

Jobs-to-be-Done

1. When a trade document arrives from SU, I want the agent to check it against the customer's requirements automatically, so I don't have to manually cross-reference every field myself.
2. When discrepancies are found, I want a ready-to-send amendment email with the exact field-level issues, so I can review and send in seconds instead of typing it out from scratch.

The Flow (human touchpoints marked in bold)

Step	Who
1. SU emails trade doc attachment	HUMAN -- SU
2. Agent extracts fields (vision model, per-field confidence)	Automated
3. Agent compares extracted fields vs. customer rules + original booking	Automated
4. Agent flags each field: match / mismatch / uncertain -- never silently approves low-confidence fields	Automated
5. Agent drafts reply (approval or amendment, field-by-field)	Automated
6. CG reviews verification result + edits draft if needed	HUMAN -- CG
7. CG clicks Send -- agent never sends on its own	HUMAN -- CG
8. SU receives reply; corrects and resubmits if flagged	HUMAN -- SU

North-Star Metric

Median time from SU document received to CG-approved reply sent. A single number that captures speed directly; accuracy is protected structurally rather than traded off against it -- uncertain fields can never be silently marked as a match, so a faster cycle time can't come from the agent cutting corners on confidence.

Failure Mode & Mitigation

Worst case: the agent marks a genuinely wrong field as a confident match, CG's reply approves bad documents, and they reach the customer -- causing customs delays or a contract penalty. **Mitigation:** the extraction step is required to report a confidence score per field; anything below threshold is forced to "uncertain" regardless of whether the text happens to match, so it can never be silently auto-approved. The agent only ever drafts -- CG's review of the field-by-field table is the final gate before anything reaches SU or the customer.